



Driving Innovation in OpenSearch, the Open Source Approach

Anandhi Bumstead

Director of Engineering at OpenSearch,
Amazon Web Services

What is the OpenSearch Project?



- Search and Analytics suite
 - Open source and community-driven
 - Now in version 2.13.0 (released April 2)
-
- **OpenSearch**
 - **OpenSearch Dashboards**
 - **Data Prepper**

How are people using OpenSearch?



Search



Observability



Security Analytics



Visualization



Machine Learning

OpenSearch Project growth

- **>500MM** Project downloads since launch
- **17 releases** 100s of new features
- **Top 4 search engine** Per DB-Engines ranking
- **70+ Partners** And growing
- **Multiple service providers** AWS, Oracle, Aiven, Bonsai-GCP
- **Public Slack channel** 2,100+ members, 100K+ messages
- **Active user forum** 5,700 members, 300K views/month

Community contributions are growing

- 2000+ Contributors outside of AWS
- *Oracle, Aiven, Aryn, ByteDance, SAP, Intel, IBM, and more*
- Weekly average of 100+ contributions



Moving toward a more open OpenSearch

- Open governance model
- Fostering more contribution
- Elevating more maintainers



Performance Improvements Since Launch

Indexing

25%

throughput
increase

Query

15-98%

latency
reduction

Storage

15-30%

disk space
reduction

Range queries

50-70%

latency
reduction

Zstandard compression

15-30%

storage
reduction

Full-text queries

15-65%

latency
reduction

Match_all queries

85%

latency
reduction

Vector search

30%

latency
reduction

Remote-backed storage

Durability

increase

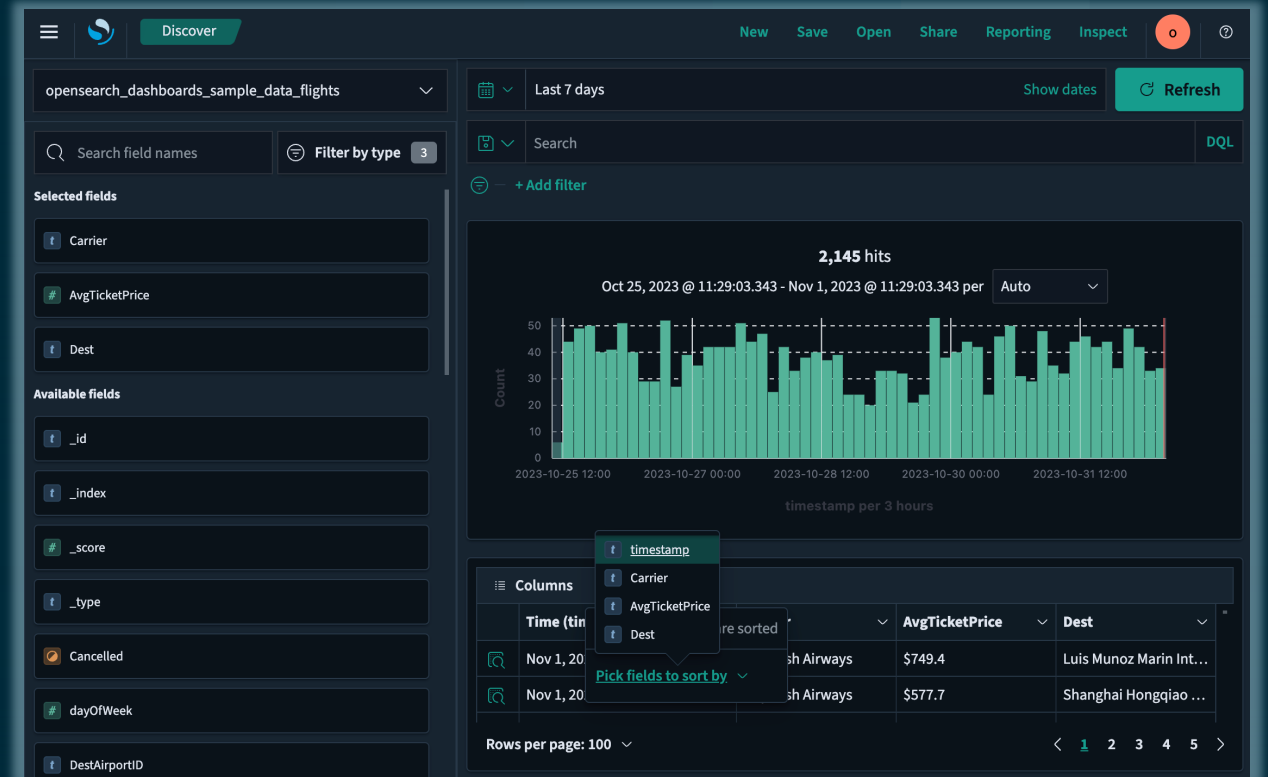
Segment replication

25%

throughput
increase

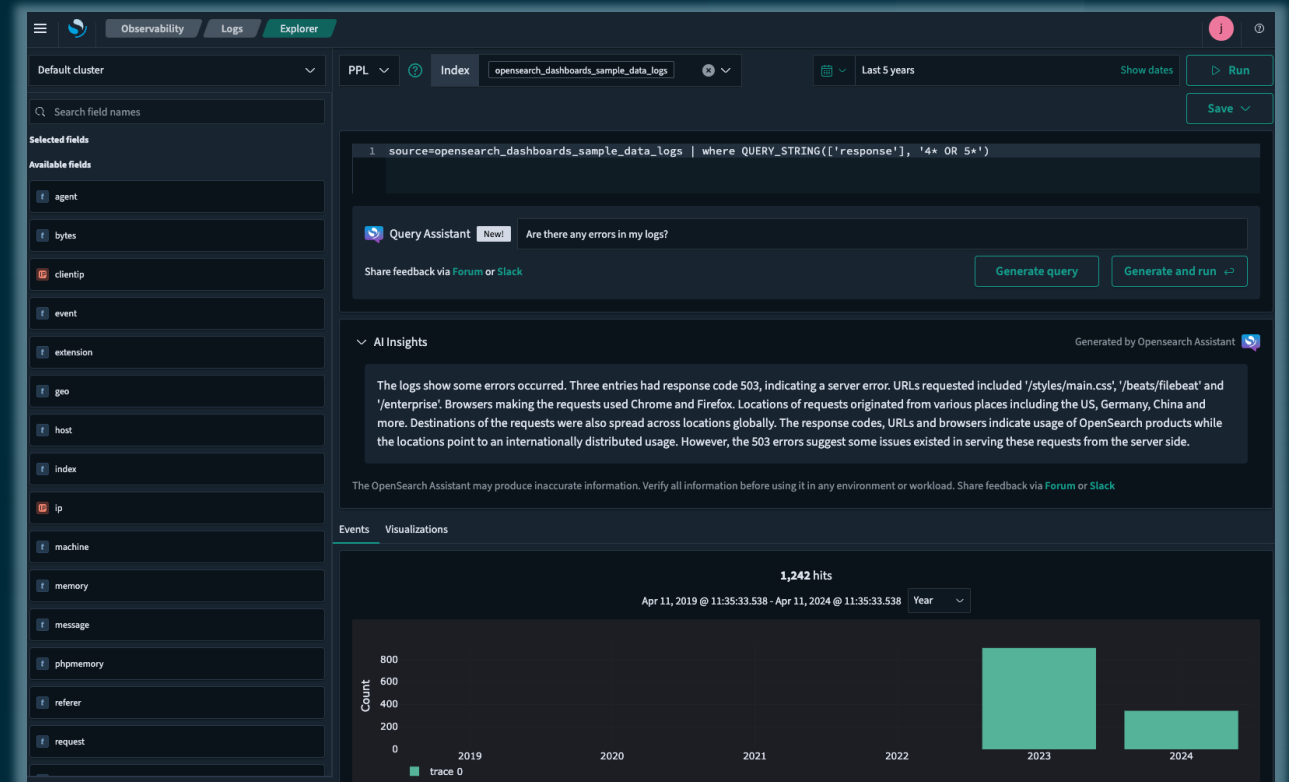
Advancing observability and analytics

- Apache Spark data
- Support for multiple data sources
- Easy-to-deploy integrations



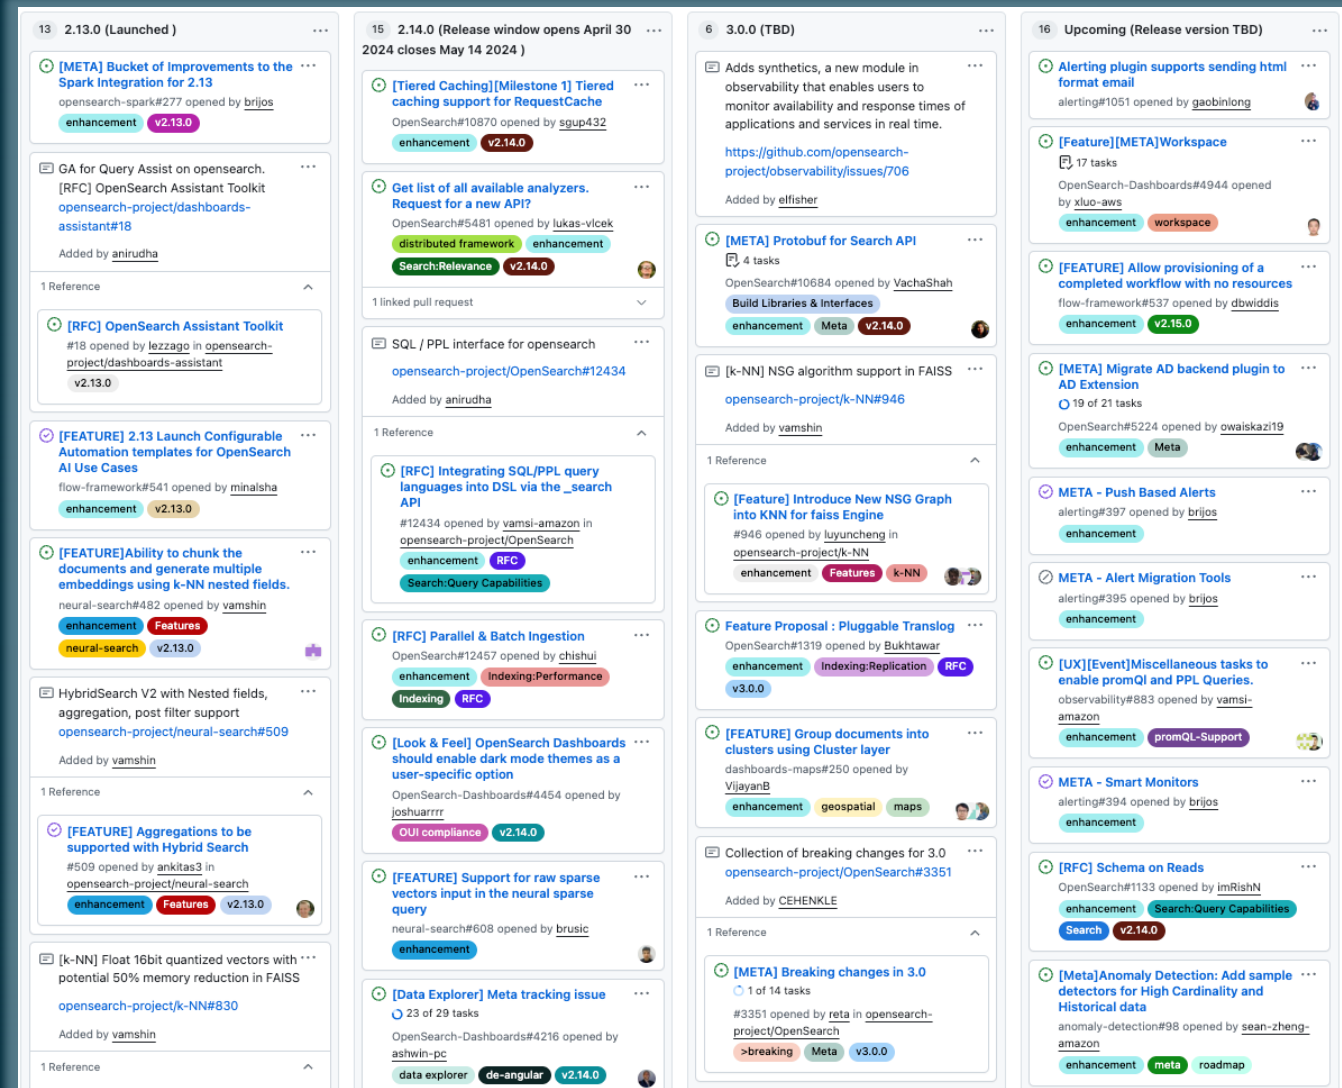
Search and machine learning innovation

- The OpenSearch Assistant Toolkit
- Flow Framework
- Enhanced hybrid search
- Improved vector quantization
- Support for pre-trained ML models



A pipeline of innovation

- Project roadmap at <https://github.com/orgs/opensearch-project/projects/1>



Get involved with OpenSearch

Multiple ways to participate

Forum - forum.opensearch.org

Slack - <https://opensearch.org/slack.html>

Events - opensearch.org/events/

Blog posts - opensearch.org/blog

Documentation - opensearch.org/docs

Code - github.com/opensearch-project/

Come join the community

Contribute and become a maintainer

Help support your favorite OS distribution

Add to the ML Commons algorithm library

Integrate OpenSearch with your favorite tools

Share your projects on the community project page

Join a user group (or create one in your city!)



Many thanks to
you, the
community!

Visit us in Booth P2!