

# Rediscover your keyword search: Expand, Enrich and Rewrite

Praveen Mohan Prasad

Analytics Specialist TAM  
AWS

Hajer Bouafif

Analytics Solutions Architect  
AWS

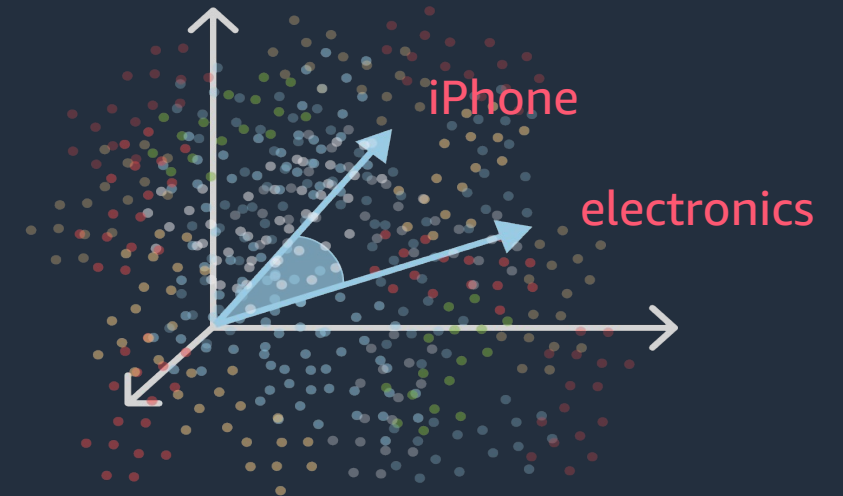
# Information Retrieval

## Sparse Lexical Search

	cat	dog	is	it	my	not	old	wolf
It is dog	0	0.25	0.22	0.25	0	0		0
my cat is old	0.3	0	0.22	0	0.3	0	0.3	0
It is not dog, it is wolf	0	0.11	0.19	0.22	0	0.13	0	0.13

Algorithm: TF-IDF, BM25  
Statistic: Frequency

## Dense Vector Search



Algorithm: kNN, ANN  
Measure: Cosine, Euclidean distance

# Strengths

## Lexical Search

- + Exact Matching
- + Interpretability
- + Less memory and Fast retrieval

## Vector Search

- + Context matching
- + Natural language understanding: RAG
- + Multimodal search

# Benchmarking on Generalisation



**Beir**  
Benchmarking IR  
18 multi-domain datasets

Beir: A heterogeneous benchmark for zero-shot evaluation of information retrieval models. [arXiv preprint arXiv:2104.08663](https://arxiv.org/abs/2104.08663).

- *Vector Search poorly generalizes on Out Of Domain (OOD) Data"*
- *Lexical search > Vector Search @ OOD*

Fine-tuning Vector Search **is complex!**

1. Training dataset
2. Data Science Expertise
3. retrain model and re-indexing

# Can we do semantic search with lexical search?

# Lexical search: limited semantic capabilities

**Document:** Exercising regularly makes **body** and **mind** **stronger** and **healthy**

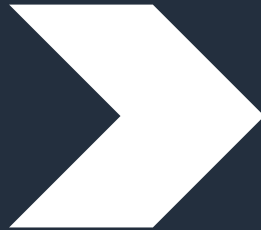
**Query:** How to **strengthen** the **physical** and **mental** **wellness**?

➡ **✗ No match**

## Challenges

Vocabulary Mismatch

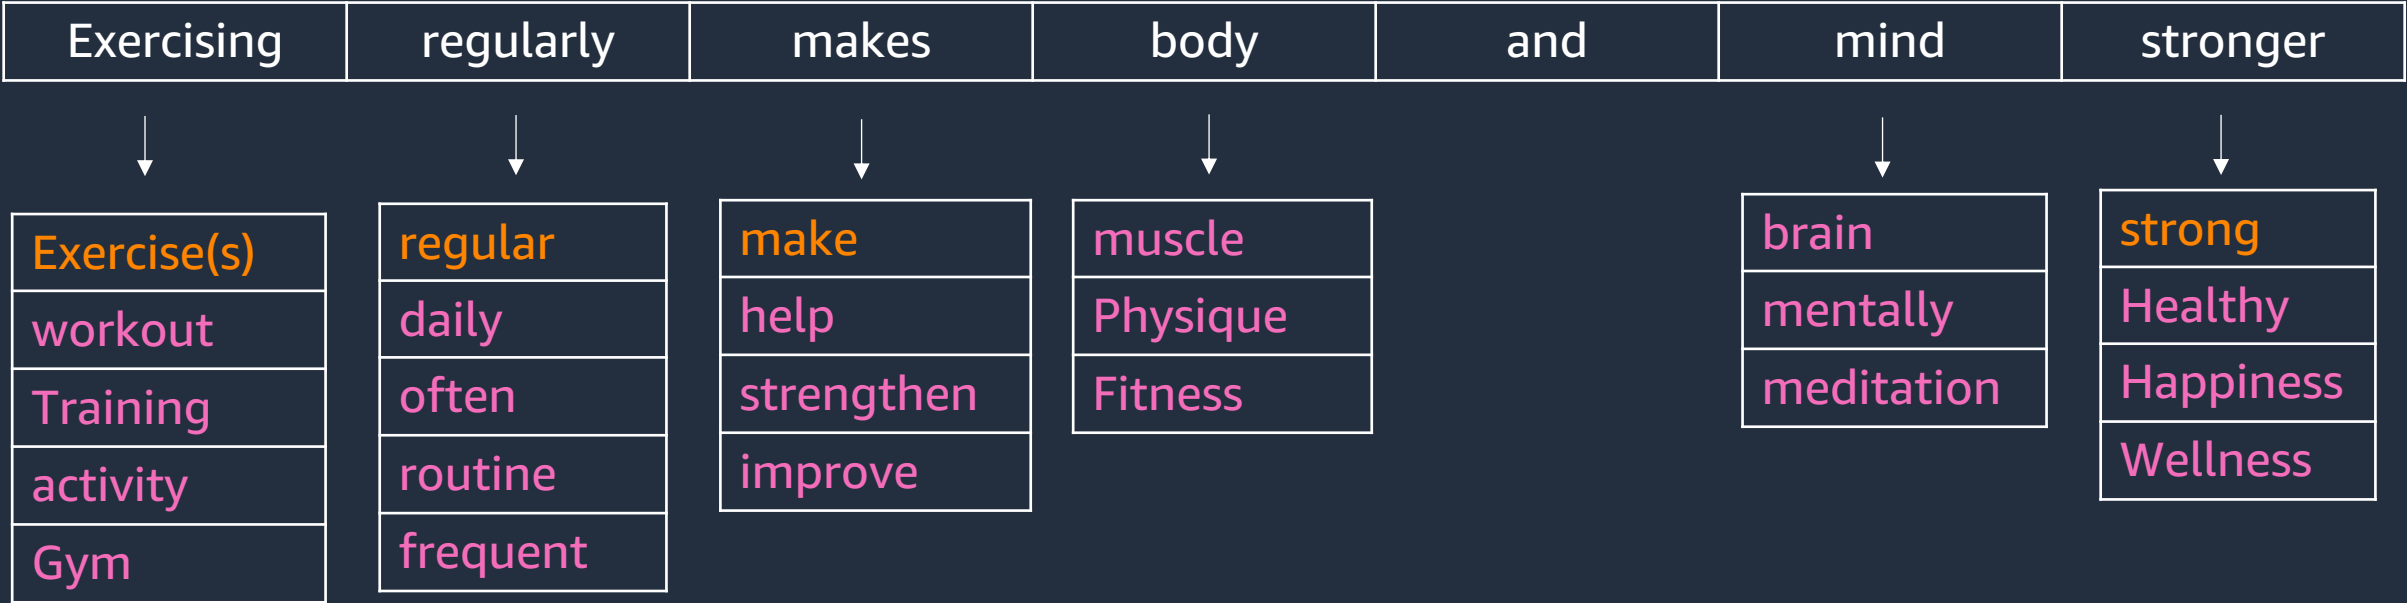
Poor semantic understanding



## Existing Solutions

- ✓ Synonyms
- ✓ Analyzers and Boosting techniques
- ✓ Learning features from User signals
- ✓ Re-ranking

# Lexical Search + Text expansion = Semantic Search



# Boost and decay the features

Document 1: "Apple Products are expensive"

```
{ 'apple': 3.09,
  'expensive': 2.58,
  'apples': 2.31,
  'cost': 2.06,
  'products': 1.78,
  'cheap': 1.61,
  'product': 1.57,
  'price': 1.28,
  'expense': 1.04,
  'best': 0.66,
  'brand': 0.46,
  'stock': 0.44,
  'chip': 0.44,
  'store': 0.25,
  'computer': 0.24,
  'offer': 0.22,
  'money': 0.21,
  'budget': 0.21,
  'good': 0.2,
  'buy': 0.19,
  'affordable': 0.19,
  'popular': 0.16,
  'gift': 0.14,
  'manufacturer': 0.12,
  'purchase': 0.09,
  'iphone': 0.09,
  'happiness': 0.06,
  'steve': 0.03,
  'amazon': 0.03,
  'hardware': 0.01}
```

Document 2: "An apple a day keeps doctor away"

```
{ 'apple': 2.45,
  'doctor': 2.26,
  'day': 2.08,
  'away': 2.05,
  'apples': 1.87,
  'doctors': 1.86,
  'daily': 1.83,
  'keep': 1.59,
  'medical': 1.19,
  'prevent': 0.97,
  'every': 0.95,
  'a': 0.94,
  'fruit': 0.91,
  'dr': 0.7,
  'keeping': 0.62,
  'help': 0.56,
  'stay': 0.46,
  'an': 0.42,
  'remove': 0.41,
  'remedy': 0.37,
  'drink': 0.35,
  'one': 0.31,
  'candy': 0.31,
  'pill': 0.31,
  'keeps': 0.26,
  'diet': 0.24,
  'eat': 0.23,
```

Query: "apple headphones"

```
{'apple': 3.298149,
 'gift': 0.22470483,
 '##phone': 1.6386933,
 'electronics': 0.18978065,
 '##phones': 2.0765076,
 'music': 0.725811,
 'dj': 0.36222592,
 'sound': 0.77559084}
```

$score(D1, Q) = 7.4521604 > score(D2, Q) = 5.5605326$

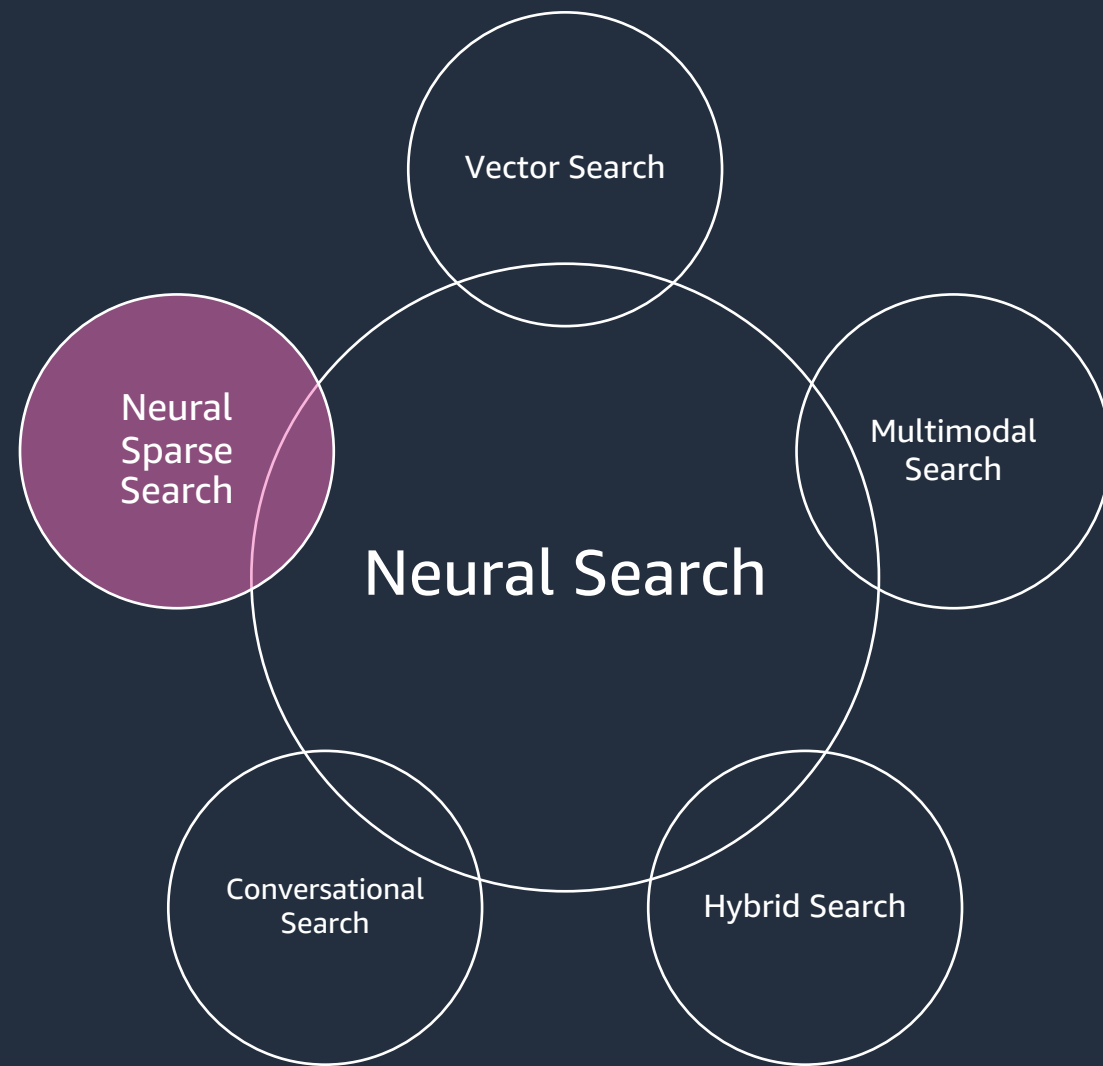
Scoring:

$$score(D, Q) = \sum w_j \cdot w_i$$





# Neural Sparse Search 2.11




# Open-source Sparse encoding models

## DOCUMENT AND SEARCH QUERY SPARSE ENCODING



Sparse encoding model

 [opensearch-project/opensearch-neural-sparse-encoding-v1](#)


 [naver/splade-v3](#)

## DOCUMENT ONLY SPARSE ENCODING



Sparse encoding model

 [opensearch-project/opensearch-neural-sparse-encoding-doc-v1](#)

 [naver/splade-v3-doc](#)

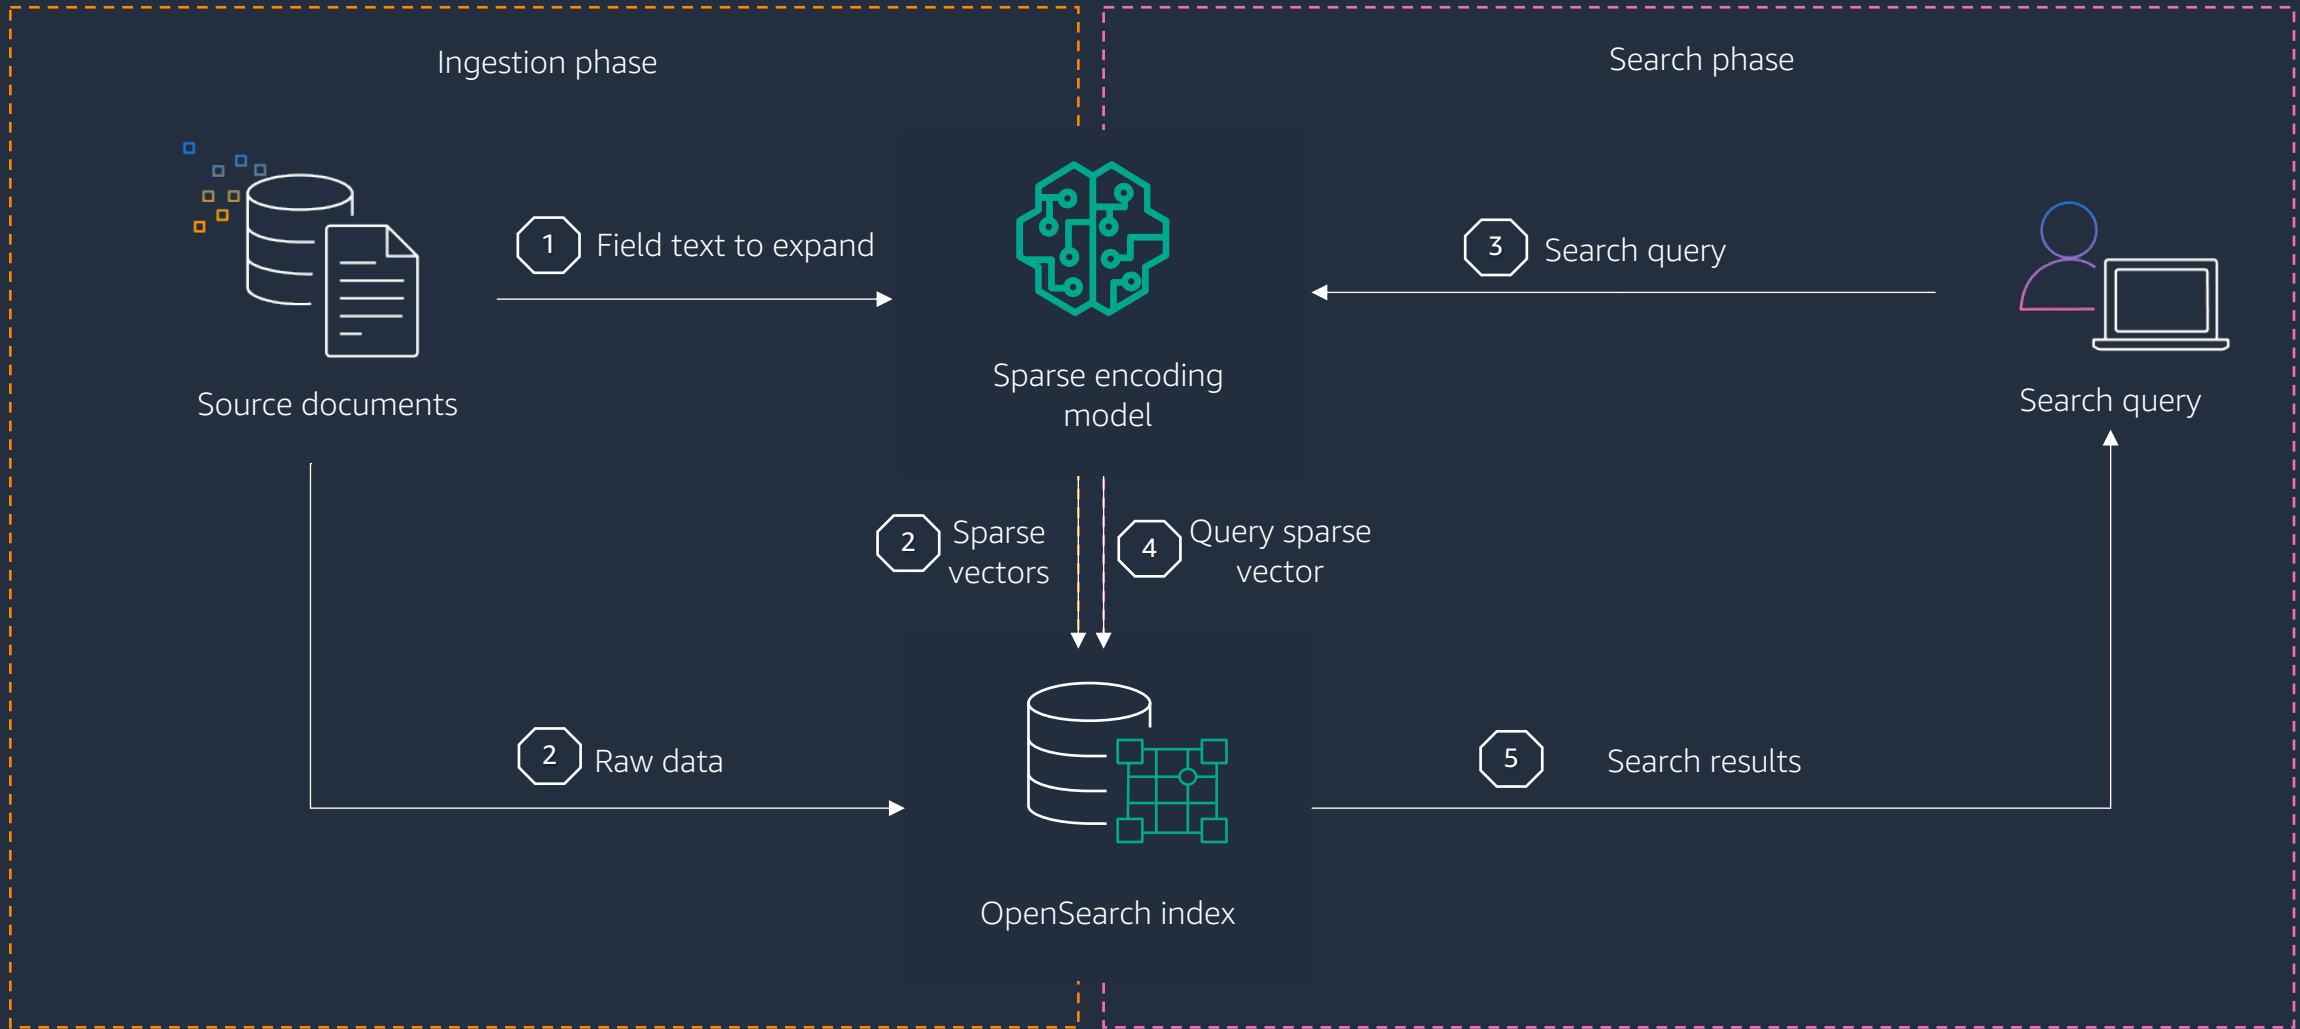
Tokenizer

[amazon/neural-sparse/opensearch-neural-sparse-tokenizer-v1](#)



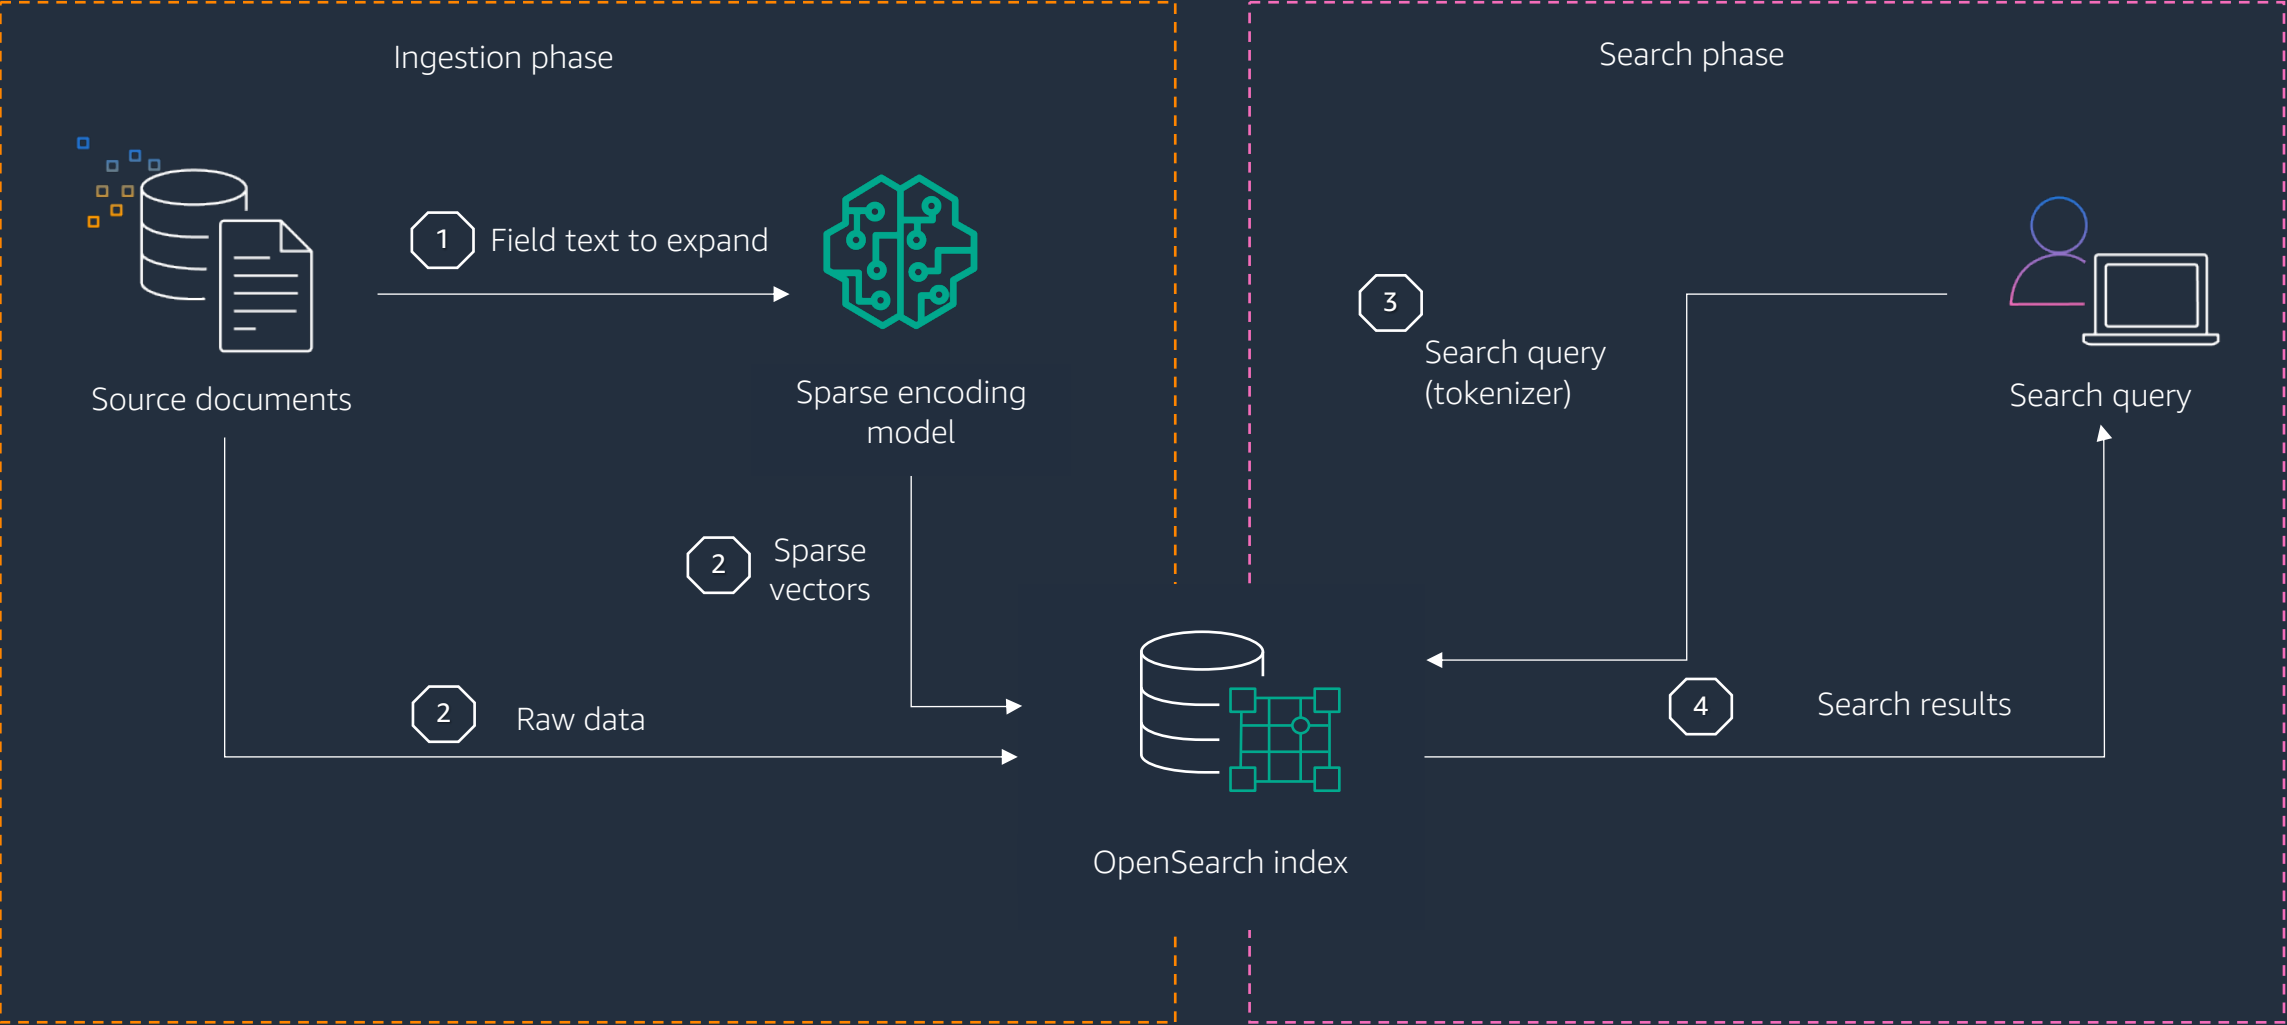
# Sparse search - High level Search architecture

## DOCUMENT AND SEARCH QUERY SPARSE ENCODING

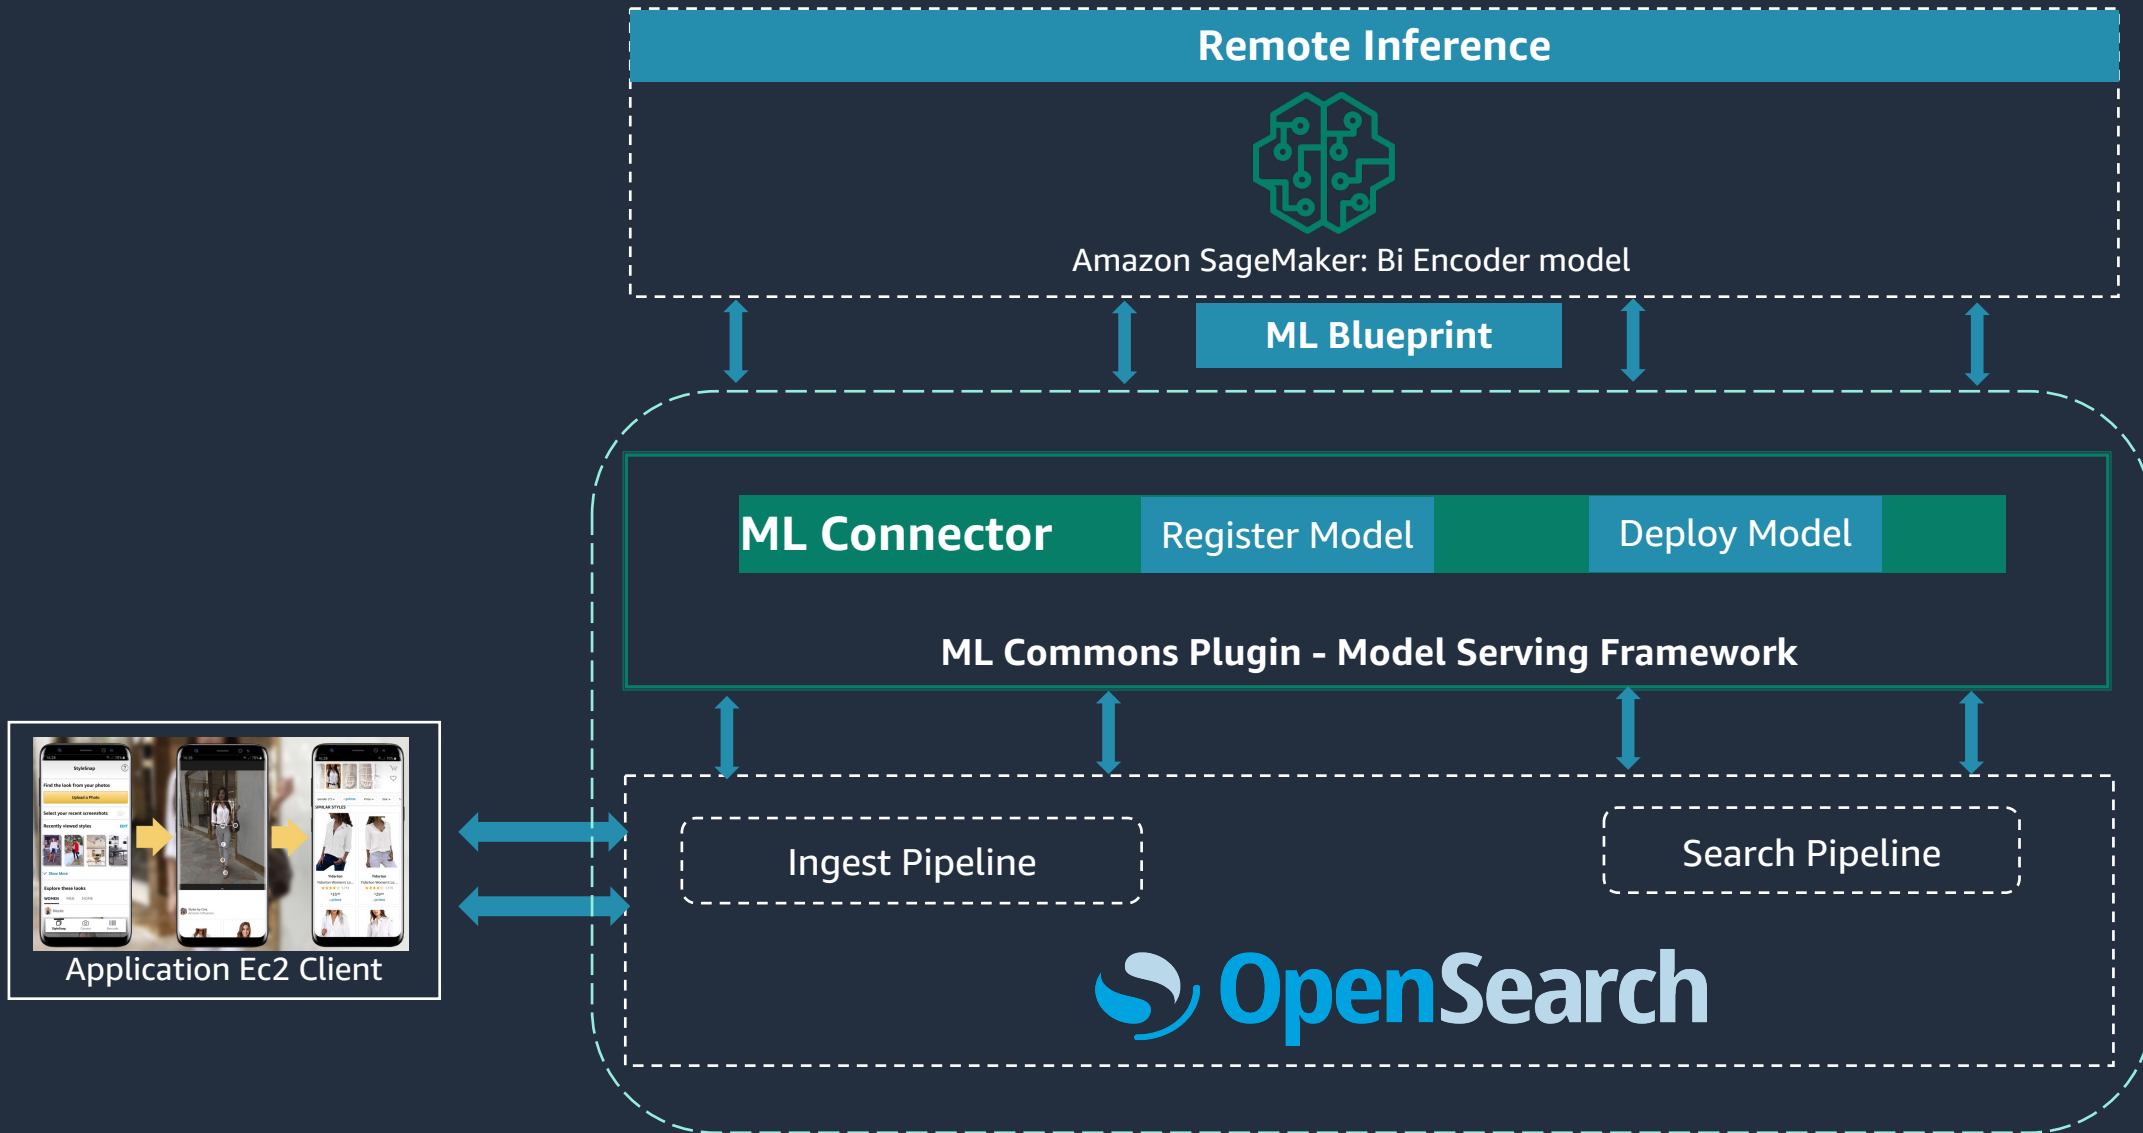


# Sparse search - High level Search architecture

## DOCUMENT ONLY SPARSE ENCODING



# Neural Sparse Search with OpenSearch



# Build Neural Sparse Search with OpenSearch

Create sparse ingest pipeline



Build the sparse index



Run neural sparse search

```
PUT /ingest/pipeline/sparse-embedding-pipeline
{
  "description": "A sparse encoding ingest
pipeline",
  "processors": [
    {
      "sparse_encoding": {
        "model_id": "4ynZg4wB33bA6yQYW-",
        "field_map": {
          "caption": "caption_sv",
          "description": "desc_sv"
        }
      }
    }
  ]
}
```

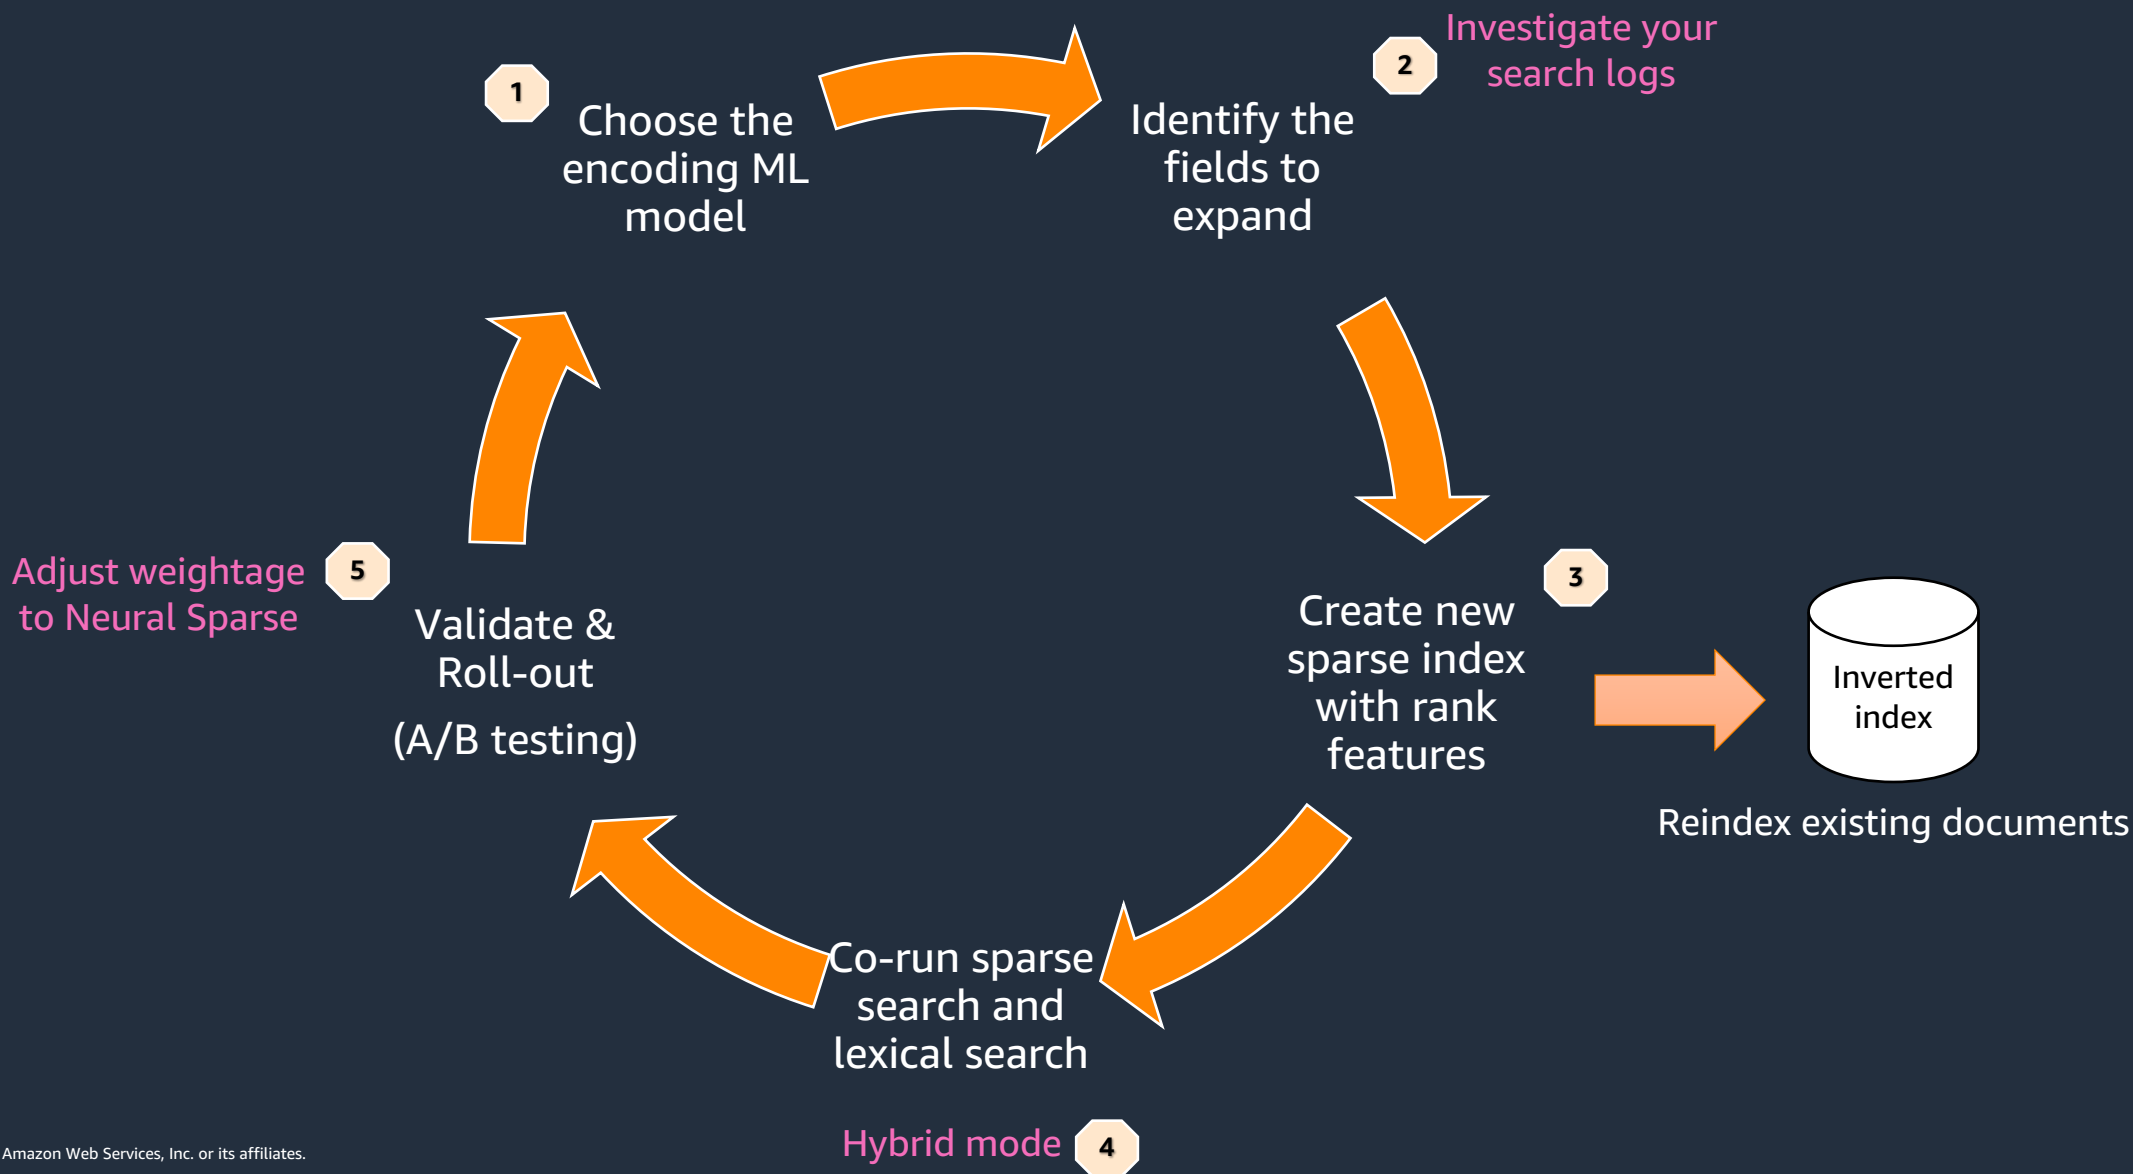
```
PUT /retail-sparse-index
{
  "settings": {
    "default_pipeline": "sparse-embedding-pipeline"
  },
  "mappings": {
    "properties": {
      "id": {
        "type": "text"
      },
      "caption": {
        "type": "text"
      },
      "caption_sv": {
        "type": "rank_features"
      },
      "description": {
        "type": "text"
      },
      "desc_sv": {
        "type": "rank_features"
      }
    }
  }
}
```

```
GET retail-sparse-index/_search
{
  "query": {
    "neural_sparse": {
      "desc_sv": {
        "query_text": "pink backpack",
        "model_id": "4ynZg4wB33bA6yQYW-"
      }
    }
  }
}
```

# Demo



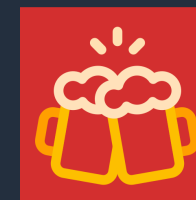
# Path to Neural Sparse search in OpenSearch





# Neural Sparse Search Benchmarks

	BM25		Dense (with TAS-B model)		Hybrid (Dense + BM25)		Neural sparse search bi-encoder		Neural sparse search document-only	
Dataset	NDCG	Rank	NDCG	Rank	NDCG	Rank	NDCG	Rank	NDCG	Rank
Trec-Covid	0.688	4	0.481	5	0.698	3	0.771	1	0.707	2
NFCorpus	0.327	4	0.319	5	0.335	3	0.36	1	0.352	2
NQ	0.326	5	0.463	3	0.418	4	0.553	1	0.521	2
HotpotQA	0.602	4	0.579	5	0.636	3	0.697	1	0.677	2
FiQA	0.254	5	0.3	4	0.322	3	0.376	1	0.344	2
ArguAna	0.472	2	0.427	4	0.378	5	0.508	1	0.461	3
Touche	0.347	1	0.162	5	0.313	2	0.278	4	0.294	3
DBPedia	0.287	5	0.383	4	0.387	3	0.447	1	0.412	2
SciDocs	0.165	2	0.149	5	0.174	1	0.164	3	0.154	4
FEVER	0.649	5	0.697	4	0.77	2	0.821	1	0.743	3
Climate FEVER	0.186	5	0.228	3	0.251	2	0.263	1	0.202	4
SciFact	0.69	3	0.643	5	0.672	4	0.723	1	0.716	2
Quora	0.789	4	0.835	3	0.864	1	0.856	2	0.788	5
Amazon ESCI	0.081	3	0.071	5	0.086	2	0.077	4	0.095	1
Average	0.419	3.71	0.41	4.29	0.45	2.71	0.492	1.64	0.462	2.64



**NDCG(Neural Sparse) > NDCG(lexical, vector)**  
By at least 5 Points

**No Fine-tuning** of models

# Latency and Memory

## Configuration:

3 data nodes – 256 GB memory, 32 vCPU  
1 ML node – 384 GB memory, 48 vCPU  
20 query clients  
MS MARCO Dataset – 8.8M docs

- Sparse(Bi-encoder) > Dense
- Lexical Search ~ Sparse(Doc encoder)

- Dense > Sparse (Bi, doc encoder) by 7.9 %
- Lexical ~ Sparse(Bi, Doc encoder)

Latency	BM25	Dense (with TAS-B model)	Neural sparse search bi-encoder	Neural sparse search document-only
P50 latency (ms)	8 ms	56.6 ms	176.3 ms	10.2ms
P90 latency (ms)	12.4 ms	71.12 ms	267.3 ms	15.2ms
P99 Latency (ms)	18.9 ms	86.8 ms	383.5 ms	22ms
Max throughput (op/s)	2215.8 op/s	318.5 op/s	107.4 op/s	1797.9 op/s
Mean throughput (op/s)	2214.6 op/s	298.2 op/s	106.3 op/s	1790.2 op/s

Memory	BM25	Dense (with TAS-B model)	Neural sparse search bi-encoder	Neural sparse search document-only
Index size	1 GB	65.4 GB	4.7 GB	6.8 GB
RAM usage	480.74 GB	675.36 GB	480.64 GB	494.25 GB
Runtime RAM delta	+0.01 GB	+53.34 GB	+0.06 GB	+0.03 GB

# Enrich and Re-write

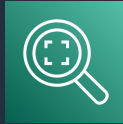


# Enrich your documents

## Images



```
{  
  "description": "Rugged Brown Leather Boots",  
  "price": "50"  
}
```

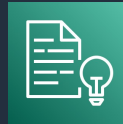


Amazon Rekognition  
or  
Object detection models

```
{  
  "description": "Rugged Brown Leather Boots",  
  "price": "50"  
  "color": "brown",  
  "category": "Apparel and Accessories",  
  "objects": "Footwear, Boot, Shoe, Clothing"  
}
```

## Text

```
{  
  "message": "Bonjour"  
}
```



Amazon Comprehend  
or  
NER models

```
{  
  "Text": "Bob ordered two  
sandwiches and three ice  
cream cones today from a  
store in Seattle."  
}
```

```
{  
  "language": "FR",  
  "message": "Bonjour"  
}  
  
{  
  {"Text": "Bob", "Type": "PERSON"},  
  {"Text": "two sandwiches", "Type": "QUANTITY"},  
  {"Text": "three ice cream cones", "Type": "QUANTITY"},  
  {"Text": "today", "Type": "DATE"},  
  {"Text": "Seattle", "Type": "LOCATION"}  
}
```

# Filters in the Query

Query: "Brown leather shoes for men under 50\$"



```
GET /_search
{
  "query": {
    "match": {
      "description": "Brown leather shoes for men"
    }
  }
}
```

```
GET _search
{
  "query": {
    "query_string": {
      "query": "Brown leather shoes for men under 50"
    }
  }
}
```

```
GET /_search
{
  "query": {
    "bool": {
      "must": [
        { "match": { "description": "shoes" } }
      ],
      "should": [
        { "term": { "color": "brown" } },
        { "term": { "gender": "male" } },
        { "term": { "material": "leather" } },
        { "range": { "price": { "lt": 50 } } }
      ]
    }
  }
}
```

# DSL Query re-writing by applying filters

TAME LLM TO EXTRACT FILTERS

Query: Brown shoes for men under \$50

Structured Query

OpenSearch Query DSL



Your goal is to structure the user's query to match the request schema provided below ....

<<query>>  
<<Schema>>  
<<examples>>



```
{  
  "query": "shoes",  
  "filter": "and"  
    (  
      eq("category", "footwear"),  
      lt("price", 50),  
      eq("gender", "male")  
      eq("color", "brown")  
    )"  
}
```



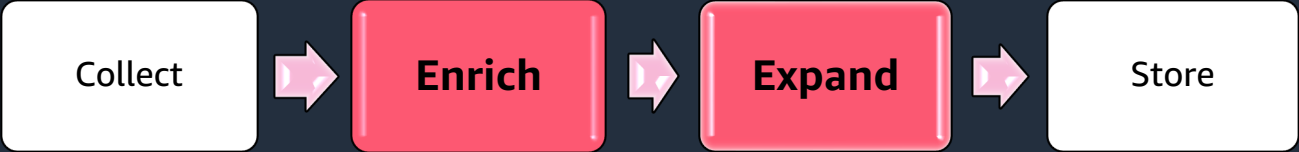
```
{  
  "query": {  
    "bool": {  
      "must": [  
        { "match": { "description": "shoes" } }  
      ],  
      "filter": [  
        { "term": { "category": "footwear" } },  
        { "range": { "price": { "lt": 50 } } },  
        { "term": { "gender": "male" } },  
        { "term": { "color": "brown" } }  
      ]  
    }  
  }  
}
```

# Demo



# (Semantic) lexical search lifecycle

Offline: **Ingestion**



Online: **Search**



# Key takeaways

1. Provide semantic search capabilities with Neural Sparse Search
  - Bi-Encoder and Document only encoder models
2. Use AI/ML connectors of OpenSearch to reduce the heavy lifting
3. Enrich your image and text documents
4. Use LLM to rewrite your DSL queries
  - Use existing frameworks and improve upon (LangChain, Haystack)
  - Measure, Evaluate and Iterate

# Thank you!

Hajer Bouafif  
bouhajer@amazon.fr

Praveen Mohan  
prasadnu@amazon.com